

# 浙江大学

## 本科实验报告

课程名称: 自然语言处理

姓 名: 王琛涵

学 院: 计算机科学与技术学院

系:

专 业: 计算机科学与技术

学 号: 3210103032

指导教师: 汤斯亮

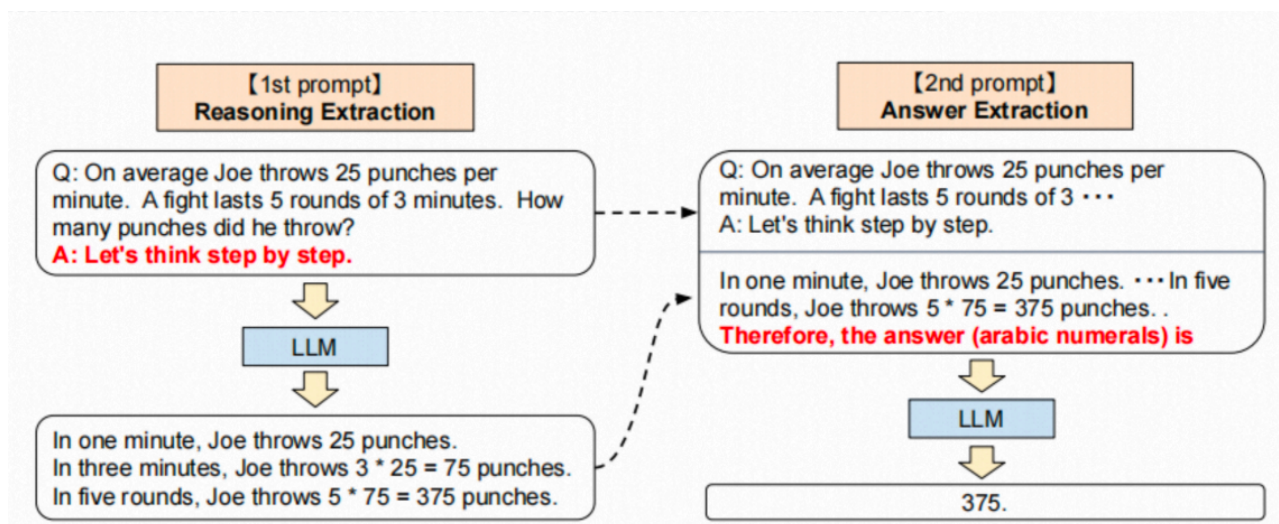
2024 年 5 月 27 日

## 1.实验原理

### 1.1 思维链

Chain of Thought (COT) 推理是指一种自然语言处理（NLP）中的推理方式，其中模型生成一系列中间步骤或“思考链”，以解决复杂的问题或任务。在传统的 NLP 任务中，模型通常被要求直接给出答案，而在 COT 推理中，模型需要展示出它是如何一步一步推理出答案的。

COT 推理特别适用于解决需要多步逻辑推理或长时间记忆的任务，如解数学问题、回答多步推理问题等。这一方法类似于人类在解决问题时先内部推理出问题的答案，再把最终答案说出来。Chain of Thought 推理通过这种方式提高了解决复杂任务的准确性，并使预训练语言模型的输出更加透明和可理解。



## 2.实验过程

### 2.1 准备步骤

#### 2.1.1 实验环境

- 环境: 华为云Notebook
- 框架: PyTorch 2.0.0, CUDA 11.7
- Python: 3.9.11

- 操作系统: Ubuntu 20.04

详细配置可以见报告末尾

### 2.1.2 模型下载

llama2

```
git clone https://gitee.com/hf-models/Llama-2-7b-chat-hf
```

llama3

```
git clone https://gitee.com/hf-models/Meta-Llama-3-8B
```

都是下载官方的，不要下中文微调版。

华为obs 不支持大于5GB的传输，所以还是用git

## 2.2 算术问题解答

实验步骤：

### 1. 定义cfg

```
cfg = edict({  
    'data_path': './cot/gsm8k_test.parquet',  
    'device_name': "cpu", # 使用CPU  
    'model_type': './Meta-Llama-3-8B',  
})
```

### 2. 数据预处理

加载数据集，提取answer中“#####”之后的部分作为正确答案。

```
# 数据处理
def extract_answer(answer):
    if '####' in answer:
        return answer.split('####')[-1].strip()
    return None

print("加载数据")
df = pd.read_parquet(cfg.data_path)
df_subset = df.head(50)
questions = df_subset['question'].tolist()
answer_data = df_subset['answer'].apply(extract_answer).tolist()
```

需要配合 *pyarrow* 对 *request* 格式进行处理

*pandas* 的 *Series* 和 *DataFrame* 是高度优化的数据结构，设计用于大规模数据处理和分析。它们带有许多额外的功能和元数据，例如索引、数据类型管理、向量化操作等。这些额外的功能在执行简单操作（如遍历或简单的数据提取）时，会带来额外的开销。所以转为 *list*，执行简单的迭代和访问操作时会更快。

### 3. 设置引导推理过程的触发句

第一步提示 *X0*：采用模板“Q:[X].A:[T]”，其中[X]是问题的槽位，[T]用于引导推理过程的触发句（默认为“Let’s think step by step”）。

```
trigger_sentence = "Let's think step by step"
first_step_prompts = ["Q: {}. A: {}".format(question,
trigger_sentence) for question in questions]
```

### 4. 加载模型

用 *pipeline* 函数来创建一个用于文本生成的管道，然后获取模型配置中的 *eos\_token\_id*，即结束标记的标识符

```
pipe = pipeline("text-generation", model=cfg.model_type,
device_map=cfg.device_name)
eos_token_id = pipe.model.config.eos_token_id
```

*pipeline* 函数会根据指定的任务类型和模型路径，自动加载模型和对应的 *tokenizer*，并配置好模型以便在指定设备上运行。

## 5.后续生成句子

通过循环遍历每个问题，使用 Llama 模型生成初步和最终答案，然后提取最终答案中的数字，与真实答案进行比较。

```
for question, real_answer in tqdm(zip(questions, answer_data),
total=len(questions)):
    print("Prompt x0:", question)
    # 调用llama模型，生成后续句子
    text = pipe(question, temperature=0.1,
eos_token_id=eos_token_id, max_new_tokens=50)
    print("Sentence z:", text[0]['generated_text'])
```

## 6.第二步提示

结合第一步提示X0与Z，构成新的提示“[X0][Z][A]”，[A]为触发模型输出答案的模板（默认为“Therefore, the answer (arabicnumerals) is”）。

```
# 在上面的循环内
# 构建完整提示
full_prompt = "{} Therefore, the answer (arabic numerals)
is".format(text[0]['generated_text'])
```

## 7.最终结果

将完成的提示作为输入送入LLaMA，获得答案预测sentence  $y^{\wedge}$ 。

```
# 在上面的循环内
# 调用llama模型，生成最终输出结果
end_text = pipe(full_prompt, eos_token_id=eos_token_id,
max_new_tokens=50)
print("Outcome:", end_text[0]['generated_text'])
```

## 8.准确率计算

对算术问题，提取预测sentence  $y^{\wedge}$ 中的数字作为预测答案，以计算模型准确率。

```

# 循环前
correct_answer_count = 0
# 循环内
    # 提取最终结果中的最后一个数字作为答案
    answer = re.findall(r'\d+', end_text[0]['generated_text'])
    predicted_answer = int(answer[-1]) if answer else None
    print("Predicted answer:", predicted_answer, 'Real answer:',
int(real_answer))
    if predicted_answer == int(real_answer):
        correct_answer_count += 1
# 计算准确率
accuracy = correct_answer_count / len(answer_data)
print("Accuracy: {:.2%}".format(accuracy))

```

## 2.3 常识问答部分

实验步骤:

1.定义cfg

2.数据预处理

将question、choices和text进行拼接

```

# 数据处理
# 将question、choices和text进行拼接
def concatenate_columns(row):
    # 构建问题字符串和选项字符串的合并
    question_part = f"{row['question']}"
    choices_part = ", ".join([f"{label}: {text}" for label, text in
zip(row['choices']['label'], row['choices']['text'])])
    return f"{question_part} {choices_part}"

print("load data")
df = pd.read_parquet(cfg.data_path)
df_subset = df.head(50)
questions = df_subset.apply(concatenate_columns, axis=1).tolist()
answer_data = df_subset['answerKey'].tolist()

```

3.第一步提示X0

采用模板“Q:[X].A:[T]”，其中[X]是问题的槽位，[T]用于引导推理过程的触发句（默认为“Let’s think step by step”）。

```
# 设置引导推理过程的触发句
trigger_sentence = "Let's think step by step. "
# 为每个问题构建第一步提示x0
promptls = ["Q: {}. A: {}".format(question, trigger_sentence) for
question in questions]
```

#### 4.调用llama模型

```
pipe = pipeline("text-generation", model=cfg.model_type,
device_map=cfg.device_name)
eos_token_id = pipe.model.config.eos_token_id
```

5.后续生成句子Z：将第一步提示发送到LLaMA中，产生后续句子z。

```
for prompt, real_answer in tqdm(zip(promptls, answer_data),
total=len(questions)):
    print("Prompt x0:", prompt)
    # 调用llama模型，生成后续句子z
    text = pipe(prompt, temperature=0.1, eos_token_id=eos_token_id)
    print("Sentence z:", text[0]['generated_text'])
```

#### 6.第二步提示

结合第一步提示X0与Z，构成新的提示“[X0][Z][A]”，[A]为触发模型输出答案的模板，此时触发句[A]格式稍有不同，应修改为“Therefore,among A through E, the answer is”。

```
# A为触发模型输出答案的模板
A = "Therefore, among A through E, the answer is"
full_prompt = "{}{}{}".format(prompt,text[0]['generated_text'],
A)
```

7.最终结果：将完成的提示作为输入送入LLaMA，获得答案预测sentence  $y^{\wedge}$ 。

```
# 调用llama模型，生成最终输出结果
end_text = pipe(full_prompt, eos_token_id=eos_token_id)
print("Outcome:", end_text[0]['generated_text'])
```

## 8. 计算准确率

### 3.实验结果

### 3.1 算术问题解答

### 3.1.1 第一次尝试 llama3

使用llama3 + cuda 。不过这个cpu也能跑，只不过稍微慢一点（可以接收的程度）

### 正确案例分析

```

48%|██████████          | 24/50 [21:30<23:43, 54.75s/it]
Prompt: Kyle bought last year's best-selling book for $19.50. This is with a 25% discount from the original price. What was the original price of the book?
Setting `pad_token_id` to `eos_token_id`:128001 for open-end generation.
Sentence Z: Kyle bought last year's best-selling book for $19.50. This is with a 25% discount from the original price. What was the original price of the book? A) $25 B) $26 C) $27 D) $28 E) $29
Answer:
Let x be the original price.
 $0.75x = 19.5$ 
 $x = 26$ 
The answer is B.
Setting `pad_token_id` to `eos_token_id`:128001 for open-end generation.
Outcome: Kyle bought last year's best-selling book for $19.50. This is with a 25% discount from the original price. What was the original price of the book? A) $25 B) $26 C) $27 D) $28 E) $29
Answer:
Let x be the original price.
 $0.75x = 19.5$ 
 $x = 26$ 
The answer is B. Therefore, the answer (arabic numerals) is 26.
Predicted answer: 26 Real answer: 26

```

问题： 提供了折扣后的价格并询问原价。

1. 提示
2. 生成的中间答案（**Sentence Z**）：

plaintext复制代码Kyle bought last year's best-selling book for \$19.50. This is with a 25% discount from the original price. What was the original price of the book? A) \$25 B) \$26 C) \$27 D) \$28 E) \$29

Answer:

Let  $x$  be the original price.

$$0.75x = 19.5$$
$$x = 26$$

The answer is B.



生成的中间答案展示了问题和答案的完整逻辑推理过程，最终得出答案是 \$26。

### 3. 完整提示（full\_prompt）：

plaintext复制代码Kyle bought last year's best-selling book for \$19.50. This is with a 25% discount from the original price. What was the original price of the book? A) \$25 B) \$26 C) \$27 D) \$28 E) \$29

Answer:

Let  $x$  be the original price.

$$0.75x = 19.5$$
$$x = 26$$

The answer is B. Therefore, the answer (arabic numerals) is

这个完整提示是在生成的中间答案基础上，添加了触发输出最终答案的模板。

#### 4. 最终输出结果 (Outcome) :

plaintext复制代码Kyle bought last year's best-selling book for \$19.50. This is with a 25% discount from the original price. What was the original price of the book? A) \$25 B) \$26 C) \$27 D) \$28 E) \$29

Answer:

Let  $x$  be the original price.

$$0.75x = 19.5$$
$$x = 26$$

The answer is B. Therefore, the answer (arabic numerals) is 26.

最终生成的答案明确地给出了数字 26，符合预期。

```

Predicted answer: 8 Real answer: 8
98% | 49/50 [49:10<00:57, 57.26s/it]
Prompt: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?
Setting `pad_token_id` to `eos_token_id`:128001 for open-end generation.
Sentence Z: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building? A) 1 B) 2 C) 3 D) 4 E) 5
Answer:
Total number of floors = 15
Total number of units per floor = 8
Total number of units in the building =
Setting `pad_token_id` to `eos_token_id`:128001 for open-end generation.
Outcome: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building? A) 1 B) 2 C) 3 D) 4 E) 5
Answer:
Total number of floors = 15
Total number of units per floor = 8
Total number of units in the building = Therefore, the answer (arabic numerals) is 4
Answer: D
Predicted answer: 4 Real answer: 30
100% | 50/50 [50:02<00:00, 60.06s/it]
Accuracy: 16.00%

```

### 3.1.3 第二次尝试 llama2

## 换用llama2测试，结果很差

```
98%|███████████████████████████████| 49/50 [02:26<00:02, 2.96s/it]
Prompt X0: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step
Sentence z: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step. We know that Richard lives in an apartment building with 15 floors. Each floor contains 8 units, so the total number of units in the building is  $15 \times 8 = 120$  units. Now,
Outcome: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step. We know that Richard lives in an apartment building with 15 floors. Each floor contains 8 units, so the total number of units in the building is  $15 \times 8 = 120$  units. Now, Therefore, the answer (arabic numerals) is 120.
Predict answer: 120
Real answer: 30
100%|████████████████████████████████████████████████████████████████████████████████| 50/50 [02:29<00:00, 2.98s/it]
Accuracy: 6.00%
```

### 3.1.3 第三次尝试 llama2

进行了max token的约束，结果变差了很多，但是推理速度却很快

```
Real answer: 0  
98%|███████████████████████| 49/50 [02:26<00:02, 2.99s/it]  
Prompt X0: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step  
Sentence z: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step. We know that Richard lives in an apartment building with 15 floors. Each floor contains 8 units, so the total number of units in the building is  $15 \times 8 = 120$  units. Now,  
Outcome: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step. We know that Richard lives in an apartment building with 15 floors. Each floor contains 8 units, so the total number of units in the building is  $15 \times 8 = 120$  units. Now,Therefore, the answer (arabic numerals) is 120.  
Predict answer: 120  
Real answer: 30  
100%|███████████████████████| 50/50 [02:29<00:00, 2.99s/it]  
Accuracy: 6.00%
```

### 3.1.3 第三次尝试 llama2

进行了max token的约束，推理速度却很快，但是结果还是很差

```
Predict answer: 8  
Real answer: 8  
98%|██████████████████████████████████████ | 49/50 [02:24<00:02, 2.98s/it]  
Prompt X0: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step  
Sentence z: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step. We know that Richard lives in an apartment building with 15 floors. Each floor contains 8 units, so the total number of units in the building is  $15 \times 8 = 120$  units. Now,  
Outcome: Q: Richard lives in an apartment building with 15 floors. Each floor contains 8 units, and 3/4 of the building is occupied. What's the total number of unoccupied units In the building?. A: Let's think step by step. We know that Richard lives in an apartment building with 15 floors. Each floor contains 8 units, so the total number of units in the building is  $15 \times 8 = 120$  units. Now, Therefore, the answer (arabic numerals) is 120.  
Predict answer: 120  
Real answer: 30  
100%|██████████████████████████████████████ | 50/50 [02:27<00:00, 2.95s/it]  
Accuracy: 10.00%
```

### 3.1.4 第四次尝试 llama2

```
# 调用llama模型，生成后续句子z
text = pipe(prompt, temperature=0.1,
eos_token_id=eos_token_id,max_new_tokens=50)
```

修改 **temperature**, 使文本多样性能力增强

**temperature = 0.3**

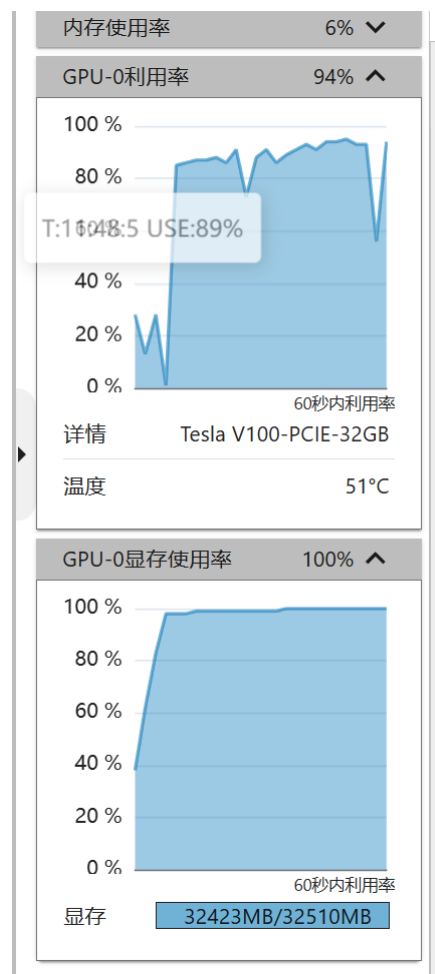


```
Predict answer: 96
Real answer: 30
100%| 50/50 [02:41<00:00, 3.22s/it]
Accuracy: 14.00%
```

## 3.2 常识选择题

### 3.2.1 第一次尝试 llama3

llama3 + cuda，显存32G差一点，尝试清理无用显存，依然不够用。




转尝试cpu，速度非常慢，一轮需要10 min +，放弃

### 3.2.2 第二次尝试 llama2

llama2比llama3小一些，cuda可以加载。





```
Real answer: E  
96%|██████████████████████████████████████ | 48/50 [09:16<00:22, 11.35s/it]  
Prompt X0: Q: The sensor would just the distance then set off an alarm, the installation expert explained it was called a what kind of sensor? A: near, B: closeness, C: here, D: proximity, E: this. A: Let's think step by step.  
Sentence z: Q: The sensor would just the distance then set off an alarm, the installation expert explained it was called a what kind of sensor? A: near, B: closeness, C: here, D: proximity, E: this. A: Let's think step by step.  
1. What is the purpose of the sensor?  
Outcome: Q: The sensor would just the distance then set off an alarm, the installation expert explained it was called a what kind of sensor? A: near, B: closeness, C: here, D: proximity, E: this. A: Let's think step by step. 1. What is the purpose of the sensor? Therefore, among A through E, the answer is (D) proximity sensor. The proximity sensor detects the distance between the object and the sensor and sets off an alarm when the distance is within a predetermined range.  
Predict answer: W  
Real answer: D  
98%|██████████████████████████████████████ | 49/50 [09:19<00:08, 8.68s/it]  
Prompt X0: Q: The man was eating lunch, but rushed when he looked at his watch, why did he rush? A: gain weight, B: late for work, C: heartburn, D: bad breath, E: early for work. A: Let's think step by step.  
Sentence z: Q: The man was eating lunch, but rushed when he looked at his watch, why did he rush? A: gain weight, B: late for work, C: heartburn, D: bad breath, E: early for work. A: Let's think step by step. 1. What was the man doing? 2  
Outcome: Q: The man was eating lunch, but rushed when he looked at his watch , why did he rush? A: gain weight, B: late for work, C: heartburn, D: bad breath, E: early for work. A: Let's think step by step. 1. What was the man doing? 2Therefore, among A through E, the answer is (B) late for work.  
Predict answer: W  
Real answer: B  
100%|██████████████████████████████████████ | 50/50 [09:20<00:00, 11.20s/it]  
Accuracy: 4.00%
```

这一次是找生成文本中最后的字母

```
answer = re.findall(r'[A-Z]', text[0][
    'generated_text'])
```

效果不好，4%的准确率基本就是随机结果

### 3.2.3 第三次尝试 llama2

约束结果，在生成结果后进行后处理，确保答案在 A-E 范围内。

```
answer = re.findall(r'[A-E]', text[0]['generated_text'])
```

```
Real answer: C
26%|██████████          | 13/50 [04:16<07:20, 11.89s/it]
Prompt X0: Q: James was cooling off too quickly. He would die if he didn't
find some way to stop what? A: loss of heat, B: revenge, C: expansion, D: re
laxation, E: calm down. A: Let's think step by step.
Sentence z: Q: James was cooling off too quickly. He would die if he didn't
find some way to stop what? A: loss of heat, B: revenge, C: expansion, D: r
elaxation, E: calm down. A: Let's think step by step. James was cooling off
too quickly, so he
Outcome: Q: James was cooling off too quickly. He would die if he didn't fi
nd some way to stop what? A: loss of heat, B: revenge, C: expansion, D: rela
xation, E: calm down. A: Let's think step by step. Q: James was cooling off
too quickly. He would die if he didn't find some way to stop what? A: loss
of heat, B: revenge, C: expansion, D: relaxation, E: calm down. A: Let's thi
nk step by step. James was cooling off too quickly, so heTherefore, among A
through E, the answer is (A) loss of heat.
Predict answer: A
Real answer: A
```

Outcome: Q: James was cooling off too quickly. He would die if he didn't find some way to stop what? A: loss of heat, B: revenge, C: expansion, D: relaxation, E: calm down. A: Let's think step by step. 【分割一下】  
Q: James was cooling off too quickly. He would die if he didn't find some way to stop what? A: loss of heat, B: revenge, C: expansion, D: relaxation, E: calm down. A: Let's think step by step. James was cooling off too quickly, so heTherefore, among A through E, the answer is (A) loss of heat.

后续生成就是之前的重复，没有出现新的有用信息。

还有的生成文本无意义生成很多次







### 3.2.5 第五次尝试 llama2

仔细研究发现， `predict answer` 的选择方式过于粗糙

AnswerTherefore, among A through E, the answer is (D) television. A television requires cable to receive broadcast signals. Radioshack, substation, cabinet, and desk do not require cable for functioning.  
Predict answer: A  
Real answer: D

llama事实上成功判断了，但是由于截取选择问题，出现错误

尝试限制结果文本的生成长度 `max_new_tokens=10`，有进步，但是也不排除随机结果的可能性

```
Real answer: D
98%|███████████████████████████████████████| 49/50 [02:18<00:02, 2.91s/it]
Prompt X0: Q: The man was eating lunch, but rushed when he looked at his watch, why did he rush? A: gain weight, B: late for work, C: heartburn, D: bad breath, E: early for work. A: Let's think step by step.
Sentence z: Q: The man was eating lunch, but rushed when he looked at his watch, why did he rush? A: gain weight, B: late for work, C: heartburn, D: bad breath, E: early for work. A: Let's think step by step. 1. What was the man doing? 2. Why did he look at his watch? 3. What did he do after looking at his watch? 4. What might have caused him to rush? 5. What was the Outcome: Q: The man was eating lunch, but rushed when he looked at his watch, why did he rush? A: gain weight, B: late for work, C: heartburn, D: bad breath, E: early for work. A: Let's think step by step. Q: The man was eating lunch, but rushed when he looked at his watch, why did he rush? A: gain weight, B: late for work, C: heartburn, D: bad breath, E: early for work. A: Let's think step by step. 1. What was the man doing? 2. Why did he look at his watch? 3. What did he do after looking at his watch? 4. What might have caused him to rush? 5. What was theTherefore, among A through E, the answer is (B) late for work.
Predict answer: B
Real answer: B
100%|███████████████████████████████████████| 50/50 [02:21<00:00, 2.82s/it]
Accuracy: 42.00%
```

## 4.参考资料

参考论文: Large Language Models are Zero-Shot Reasoners

[https://github.com/mindspore-courses/step\\_into\\_llm/blob/master/Season2.step\\_into\\_llm/04.LLaMA/llama\\_infer.py](https://github.com/mindspore-courses/step_into_llm/blob/master/Season2.step_into_llm/04.LLaMA/llama_infer.py)

<https://huggingface.co/meta-llama/Meta-Llama-3-8B>

## 5. 环境配置（心得与建议）

因为默认的镜像版本都很低，而且华为不给予**sudo**权限来升级，所以只能找一个版本比较高的镜像来配置环境。

1. 找到并点击前往AI Gallery获取更多镜像

<

创建Notebook

\* 名称

描述

0/1

\* 自动停止

开启该选项后，该Notebook实例将在运行时长超出您所选择的时长后，自动停止。

1 小时

2 小时

4 小时

6 小时

自定义

\* 镜像

公共镜像

自定义镜像

?

[前往AI Gallery获取更多镜像](#)

名称	描述
<div><div><div></div></div>pytorch1.8-cuda10.2-cudnn7-ubuntu18.04</div>	CPU、GPU通用算法开发和训练基础镜像，预置AI引擎PyTorch1.8

\* 资源类型

公共资源池

专属资源池

\* 类型

CPU

GPU

## 2. 点击第一条



所有

我的收藏

所有标签

综合排序

Pytorch\_2.0.0-cuda\_11.7-py\_3.9.11-ubuntu\_20.04

推荐



ModelArts | 发布于10个月以前

2条评论

MindSpore\_2.0.0-Cuda\_11.6-py\_3.9-Ubuntu\_20.04

推荐



ModelArts | 发布于10个月以前


0条评论

PyTorch1.11.0-Cuda11.3-Ubuntu18.04



ModelArts | 发布于18个月以前

0条评论



会员计划

### 3. 复制他的url

公有云镜像

Pytorch\_2.0.0-cuda\_11.7-py\_3.9.11-ubuntu\_20.04

发布于10个月以前

☆ 37 点赞 3 1518 评论

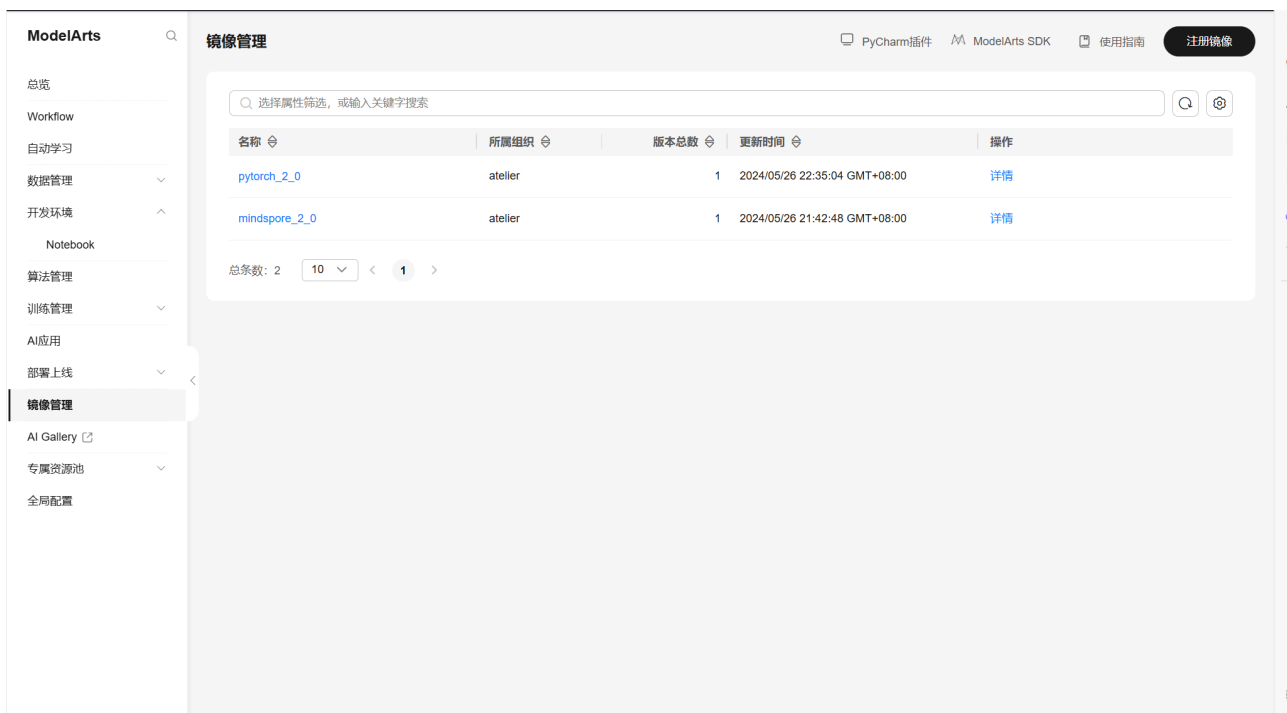
CPU and GPU general algorithm development and training, preconfigured with AI engine PyTorch2.0.0 and cuda11.7

AI引擎框架	是否使用 GPU (CUDA 版本)	URL	包含的依赖项	
			PyPI 程序包	Ubuntu 软件包
		华北-北京四 swr.cn-north-4.myhuaweicloud.com/atelier/pytorch_2_0:pytorch_2.0.0-cuda_11.7-py_3.9.11-ubuntu_20.04-x86_64-20230727142019-7d74011 华东-上海一 swr.cn-east-3.myhuaweicloud.com/atelier/pytorch_2_0:pytorch_2.0.0-cuda_11.7-py_3.9.11-ubuntu_20.04-x86_64-	torch torchvision ipykernel ipython matplotlib moxing-framework	automake build-essential ca-certificates cmake cpp curl ffmpeg g++ gcc gfortran git git-lfs grep



会员计划

### 4. 回到 modelArts， 点击注册镜像



## 5. 把复制的url粘贴于此

< | 注册镜像 ?

★ 镜像源

查看可选镜像

描述

0/256

★ 架构

X86\_64

ARM

★ 类型

☒ CPU

☐ GPU

## 6. 注册完成后，新建notebook时选择自定义镜像即可

